

CS / EE 320

Computer Organization and

Assembly Language

Spring 2025

Lecture 11

Shahid Masud

Topics: Introduction to Floating Point Numbers, Basic Operations

- Introduction to Floating Point Numbers
- Conversion from Decimal to Floating Point
- IEEE 754 Floating Point Standard Representation
- Concept of Normalization
- Concept of Biased Exponent
- Different Precision Specification in IEEE 754
- Range and Accuracy Calculations in Floating Point Standard
- Special Representations NaN, Inf, etc.

**Quiz Next
Week**

Common Physical Constants

Mostly Floating-Point Numbers

PHYSICAL CONSTANT

Speed of light	c	$3 \times 10^8 \text{ m/s}$
Planck constant	h	$6.63 \times 10^{-34} \text{ J s}$
	hc	1242 eV-nm
Gravitation constant	G	$6.67 \times 10^{-11} \text{ m}^3 \text{ kg}^{-1} \text{ s}^{-2}$
Boltzmann constant	k	$1.38 \times 10^{-23} \text{ J/K}$
Molar gas constant	R	8.314 J/(mol K)
Avogadro's number	N_A	$6.023 \times 10^{23} \text{ mol}^{-1}$
Charge of electron	e	$1.602 \times 10^{-19} \text{ C}$
Permeability of vacuum	μ_0	$4\pi \times 10^{-7} \text{ N/A}^2$
Permittivity of vacuum	ϵ_0	$8.85 \times 10^{-12} \text{ F/m}$
Coulomb constant	$\frac{1}{4\pi\epsilon_0}$	$9 \times 10^9 \text{ N m}^2/\text{C}^2$
Faraday constant	F	96485 C/mol
Mass of electron	m_e	$9.1 \times 10^{-31} \text{ kg}$
Mass of proton	m_p	$1.6726 \times 10^{-27} \text{ kg}$
Mass of neutron	m_n	$1.6749 \times 10^{-27} \text{ kg}$
Atomic mass unit	u	$1.66 \times 10^{-27} \text{ kg}$
Atomic mass unit	u	$931.49 \text{ MeV}/c^2$
Stefan-Boltzmann constant	σ	$5.67 \times 10^{-8} \text{ W/(m}^2 \text{ K}^4)$
Rydberg constant	R_∞	$1.097 \times 10^7 \text{ m}^{-1}$
Bohr magneton	μ_B	$9.27 \times 10^{-24} \text{ J/T}$
Bohr radius	a_0	$0.529 \times 10^{-10} \text{ m}$
Standard atmosphere	atm	$1.01325 \times 10^5 \text{ Pa}$
Wien displacement constant	b	$2.9 \times 10^{-3} \text{ m K}$

- Representation for non-integral numbers
 - Including very small and very large numbers

- Like scientific notation

- -2.34×10^{56}

- $+0.002 \times 100^{-4}$

- $+987.02 \times 100^9$

Normalized?

Not Normalized?

- In binary

- $\pm 1.xxxxxxx_2 \times 2^{yyyy}$

- Types `float` and `double` in C for different accuracy

Three fields in Floating Point Numbers

Sign bit	Exponent	Significand or Mantissa
----------	----------	-------------------------

- $\pm \text{.significand} \times 2^{\text{exponent}}$
- Point is actually fixed somewhere between sign bit and body of mantissa
- Exponent indicates place value (point position)

$$(-1)^S \times F \times 2^E$$

S = Sign

F = Fraction (fixed point number)

usually called **Mantissa** or **Significand**

E = Exponent (positive or negative integer)

- How to divide a word into S, F and E?
- How to represent S, F and E?

- Standard for floating-point storage
- 32 and 64 bit standards
- 8 and 11 bit exponent respectively
- Extended formats (both mantissa and exponent) for intermediate results
- Representation: sign, exponent, fraction
 - 0: 0, 0, 0
 - -0: 1, 0, 0
 - Special:
 - Plus infinity: 0, all 1s, 0
 - Minus infinity: 1, all 1s, 0
 - NaN; 0 or 1, all 1s, != 0. etc.

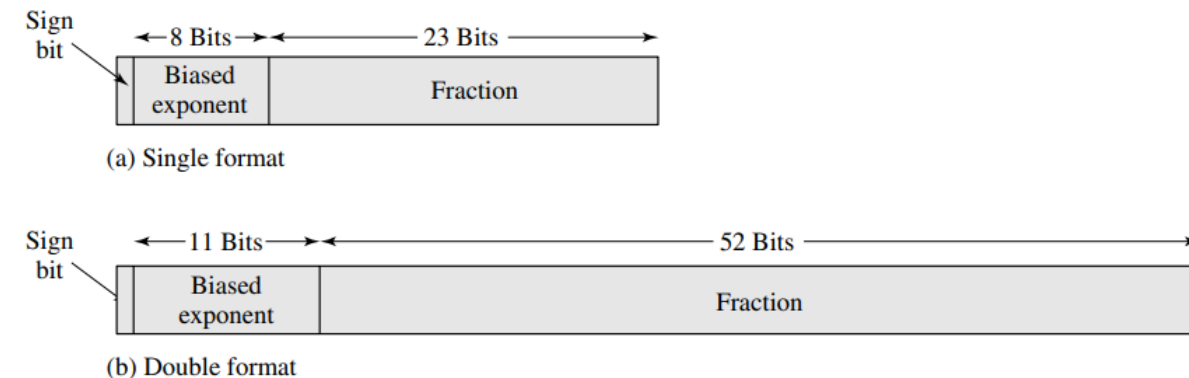


Figure 9.21 IEEE 754 Formats

- Defined by **IEEE Std 754-1985**
- Developed in response to divergence of representations
 - Portability issues for scientific code
- Now almost universally adopted
- Three representations in latest IEEE 754 standard:
 - **Half precision** (16-bit)
 - **Single precision** (32-bit)
 - **Double precision** (64-bit)

single: 8 bits
double: 11 bits

single: 23 bits
double: 52 bits



$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 \Rightarrow non-negative, 1 \Rightarrow negative)
- Normalize significand: $1.0 \leq |\text{significand}| < 2.0$
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the “1.” restored
- Exponent: excess representation: actual exponent + Bias
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1023

IEEE 754 Precision types



Name	Significand Bits	Exponent Bits	Exponent Bias
Half Precision	11	5	+15
Single Precision	24	8	+127
Double Precision	53	11	+1023
Quadruple Precision	113	15	+16383

Double precision	010000000001001001000011111101101010100010001000010110100011000
Single precision	01000000010010010000111111011011
Half precision	0100001001001000

- What number is represented by the single-precision float

11000000**1**01000...00

- $S = 1$
- Fraction = $01000...00_2$
- Exponent = $10000001_2 = 129$
- $x = (-1)^1 \times (1 + 01_2) \times 2^{(129 - 127)}$
 $= (-1) \times 1.25 \times 2^2$
 $= -5.0$

- FP numbers are usually normalized
- i.e. Exponent is adjusted so that leading bit (MSB) of mantissa is 1
- Since it is always 1 there is no need to store it
- (Scientific notation where numbers are normalized to give a single digit before the decimal point e.g. 3.123×10^3)

- Represent -0.75
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000...00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 01111111110_2$
- Single: $1011111101000...00$
- Double: $1011111111101000...00$

Some Examples of Converting to FP



- Convert given Decimal number to Floating Point Representation
- Convert to Decimal Representation from Floating Point

Some Floating-Point Conversions



(a) Format

$$\begin{aligned} 1.1010001 \times 2^{10100} &= 0 \ 10010011 \ 101000100000000000000000 = 1.6328125 \times 2^{20} \\ -1.1010001 \times 2^{10100} &= 1 \ 10010011 \ 101000100000000000000000 = -1.6328125 \times 2^{20} \\ 1.1010001 \times 2^{-10100} &= 0 \ 01101011 \ 101000100000000000000000 = 1.6328125 \times 2^{-20} \\ -1.1010001 \times 2^{-10100} &= 1 \ 01101011 \ 101000100000000000000000 = -1.6328125 \times 2^{-20} \end{aligned}$$

(b) Examples

Figure 9.18 Typical 32-Bit Floating-Point Format

Examples of Floating-Point Representation



$-(13.45)_{10}$
= $(1101.01\ 1100\ 1100\ 1100\ \dots\dots)^2$; this is un-normalized
= $(1.10101\ 1100\ 1100\ 1100\ 1100\ 1) \times 2^3$; normalized
Fraction part is 10101 1100 1100 1100 1
Biased Exponent is $3+127 = 130$
Sign = 1

5.0345
= $101.0000\ 1000\ 1101\ 0100\ 1111\ 110$; this is un-normalized
= $1.01\ 0000\ 1000\ 1101\ 0100\ 1111\ 110 \times 2^2$; normalized
Biased Exponent = $2+127 = 129 = (1000\ 0001)_2$
Fraction = 01 0000 1000 1101 0100 1111 110
Sign = 0

- Exponents **00000000** and **11111111** reserved
- Smallest value
 - Exponent: 00000001
 \Rightarrow actual exponent = $1 - 127 = -126$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
 \Rightarrow actual exponent = $254 - 127 = +127$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
 - Exponent: 000000000001
 \Rightarrow actual exponent = $1 - 1023 = -1022$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 111111111110
 \Rightarrow actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

- Relative precision
 - all fraction bits are significant
 - Single: approx 2^{-23}
 - Equivalent to $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$ decimal digits of precision
 - Double: approx 2^{-52}
 - Equivalent to $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$ decimal digits of precision

Expressible Numbers

- Negative numbers between $-(2 - 2^{-23}) \times 2^{128}$ and -2^{-127}
- Positive numbers between 2^{-127} and $(2 - 2^{-23}) \times 2^{128}$

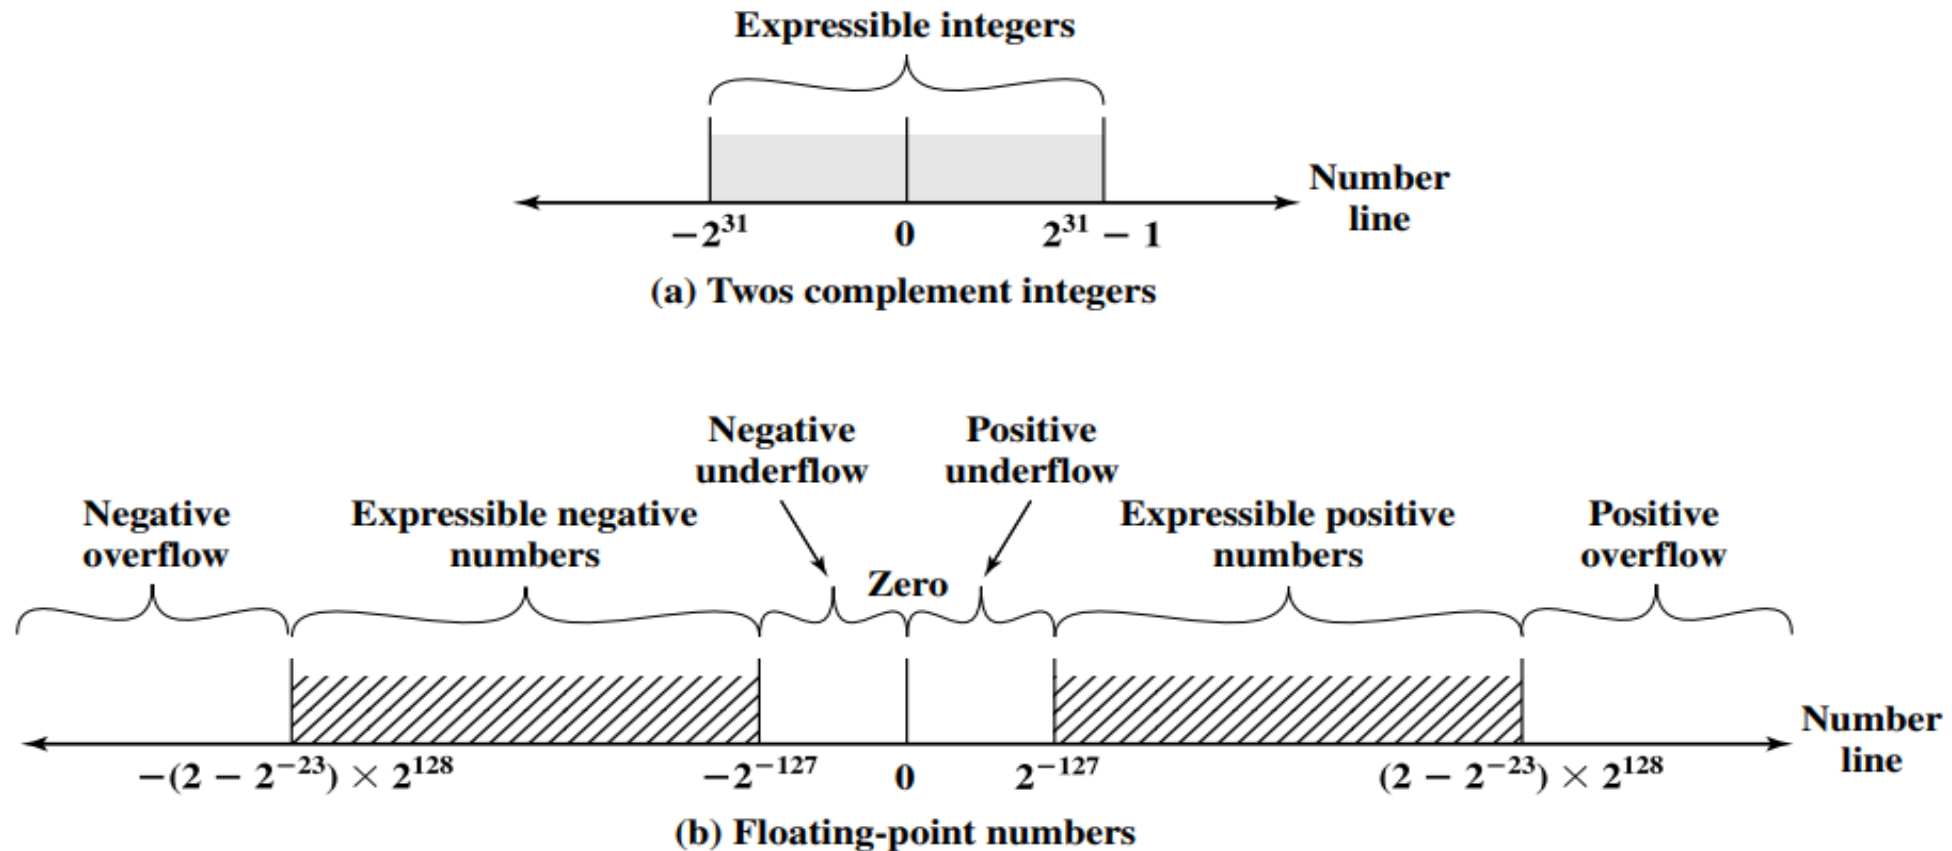


Figure 9.19 Expressible Numbers in Typical 32-Bit Formats

More about range of Floating-Point numbers



Five regions on the number line are not included in these ranges:

- Negative numbers less than $-(2 - 2^{-23}) \times 2^{128}$, called **negative overflow**
- Negative numbers greater than 2^{-127} , called **negative underflow**
- Zero
- Positive numbers less than 2^{-127} , called **positive underflow**
- Positive numbers greater than $(2 - 2^{-23}) \times 2^{128}$, called **positive overflow**

Overflow and Underflow in Floating-Point



largest positive/negative number (SP) =

$$\pm(2 - 2^{-23}) \times 2^{127} \cong \pm 2 \times 10^{38}$$

smallest positive/negative number (SP) =

$$\pm 1 \times 2^{-126} \cong \pm 2 \times 10^{-38}$$

Largest positive/negative number (DP) =


$$\pm(2 - 2^{-52}) \times 2^{1023} \cong \pm 2 \times 10^{308}$$

Smallest positive/negative number (DP) =

$$\pm 1 \times 2^{-1022} \cong \pm 2 \times 10^{-308}$$

Uneven Distribution in Floating Point

Decimal Representation:



Binary Representation:





FIGURE 2-5 Comparison of a Set of Numbers Defined by Digital and Binary Representation

Floating Point – Denormals and Distribution



$[A, B]$ — negative floating-point numbers (normalized)

$[D, E]$ — positive floating-point numbers (normalized)

(B, b) & (d, D) — denormals

C — zero


$> E$ — positive overflow

$< A$ — negative overflow

(B, C) — negative underflow (normalized)

(C, D) — positive underflow (normalized)

(a)



(b)

Special Representation – Subnormal



Zero: Zero is a special value denoted with an exponent field of 0 and a mantissa of 0.

Denormalized: If the exponent is all zeros, but the mantissa is not then the value is a *denormalized* number


Infinity: The values +infinity and -infinity are denoted with an exponent of all ones and a mantissa of all zeros. The sign bit distinguishes between negative infinity and positive infinity. Being able to denote infinity as a specific value is useful because it allows operations to continue past overflow situations. Operations with infinite values are well defined in IEEE.

Indeterminate: The value *indeterminate* is represented by an exponent of all ones, a mantissa with a leading one followed by all zeros, and a sign bit of one. This value is used to represent results that are indeterminate, such as (infinity - infinity), or (0 x infinity).


Not A Number: The value NaN (*Not a Number*) is used to represent a value that is an error of some form. This is represented with an exponent field of all ones and a zero sign bit or a mantissa that it not 1 followed by zeros.

Type	Exp	Fraction	Sign
Positive Zero	0	0	0
Negative Zero	0	0	1
Denormalised numbers	0	non zero	any
Normalised numbers	$1..2^e - 2$	any	any
Infinities	$2^e - 1$	0	any
NaN	$2^e - 1$	non zero	any

Subnormal or Denormal Numbers



(a) 32-bit format without subnormal numbers



(b) 32-bit format with subnormal numbers

- IEEE Std 754 specifies additional rounding control
 - Extra bits of precision (**guard**, **round**, **sticky**)
 - Choice of rounding modes
 - Allows programmer to fine-tune numerical behavior of a computation
- Not all FP units implement all options
 - Most programming languages and FP libraries just use defaults
- Trade-off between hardware complexity, performance, and market requirements

IEEE 754 Floating-Point Representation



Single Precision IEEE 754

Sign	Exponent	Fraction
1 bit	8 bits	23 bits

← 32 bits →

G R S

Guard Bits

For Accuracy

Double Precision IEEE 754

Sign	Exponent	Fraction
1 bit	11 bits	52 bits

← 64 bits →

G R S

Guard Bits

- Precision is lost when some bits are shifted to right of the rightmost bit or are thrown
- Three extra bits are used internally -
G (guard), R (round) and S (sticky)
 - G and R are simply the next two bits after LSB
 - S = 1 iff any bit to right of R is non-zero

1. 11010110101100010110110 GRS

ROUNDING Another detail that affects the precision of the result is the rounding policy. The result of any operation on the significands is generally stored in a longer register. When the result is put back into the floating-point format, the extra bits must be eliminated in such a way as to produce a result that is close to the exact result. This process is called **rounding**.

A number of techniques have been explored for performing rounding. In fact, the IEEE standard lists four alternative approaches:

- **Round to nearest:** The result is rounded to the nearest representable number.
- **Round toward $+\infty$:** The result is rounded up toward plus infinity.
- **Round toward $-\infty$:** The result is rounded down toward negative infinity.
- **Round toward 0:** The result is rounded toward zero.

Let us consider each of these policies in turn. **Round to nearest** is the default rounding mode listed in the standard and is defined as follows: The representable value nearest to the infinitely precise result shall be delivered.

- if $G=1$ & $R=1$, add 1 to LSB
- if $G=0$ & $R=0$ or 1, no change
- if $G=1$ & $R=0$, look at S
 - if $S=1$, add 1 to LSB
 - if $S=0$, round to the nearest “even”
i.e., add 1 to LSB if $\text{LSB} = 1$

Another view of Rounding Scheme



G Guard	R Round	S Sticky	Rounding Applied
0	0	0	Truncate
0	0	1	Truncate
0	1	0	Truncate
0	1	1	Truncate
1	0	0	Round to Even
1	0	1	Round Up
1	1	0	Round Up
1	1	1	Round Up

$S = \text{OR}(\text{All bits in S and to the right of S})$

Round to Even:

If $S = 0$, do nothing

If $S = 1$, add +1

$$N = (-1)^S \times (1 + F) \times 2^E$$

- ❖ A signed-magnitude system for the fractional part and a biased notation for the exponent
- ❖ Three subfields
 - ❖ Sign S
 - ❖ Fraction F (or Significand or Mantissa)
 - ❖ Exponent E
- ❖ Sign bit is 0 for positive numbers, 1 for negative numbers
- ❖ Fractions always start from **1**.xxxx, hence the integer **1** is not written (register has xxxx)
- ❖ Exponent is biased by +127 (add 127 to whatever is in register bits)
- ❖ Normalize: Express numbers in the standard format by shifting of bits and adding / subtracting from Exponent register

The Use of Guard Bits

$$\begin{aligned}x &= 1.000\dots00 \times 2^1 \\ \underline{-y} &= \underline{0.111\dots11} \times 2^1 \\ z &= 0.000\dots01 \times 2^1 \\ &= 1.000\dots00 \times 2^{-22}\end{aligned}$$

(a) Binary example, without guard bits

$$\begin{aligned}x &= .100000 \times 16^1 \\ \underline{-y} &= \underline{.0FFFFFF} \times 16^1 \\ z &= .000001 \times 16^1 \\ &= .100000 \times 16^{-4}\end{aligned}$$

(c) Hexadecimal example, without guard bits

$$\begin{aligned}x &= 1.000\dots00 \ 0000 \times 2^1 \\ \underline{-y} &= \underline{0.111\dots11 \ 1000} \times 2^1 \\ z &= 0.000\dots00 \ 1000 \times 2^1 \\ &= 1.000\dots00 \ 0000 \times 2^{-23}\end{aligned}$$

(b) Binary example, with guard bits

$$\begin{aligned}x &= .100000 \ 00 \times 16^1 \\ \underline{-y} &= \underline{.0FFFFFF \ F0} \times 16^1 \\ z &= .000000 \ 10 \times 16^1 \\ &= .100000 \ 00 \times 16^{-5}\end{aligned}$$

(d) Hexadecimal example, with guard bits

Floating-Point Numbers

$$X = X_S \times B^{X_E}$$

$$Y = Y_S \times B^{Y_E}$$

Arithmetic Operations

$$\left. \begin{aligned} X + Y &= (X_S \times B^{X_E - Y_E} + Y_S) \times B^{Y_E} \\ X - Y &= (X_S \times B^{X_E - Y_E} - Y_S) \times B^{Y_E} \end{aligned} \right\} X_E \leq Y_E$$

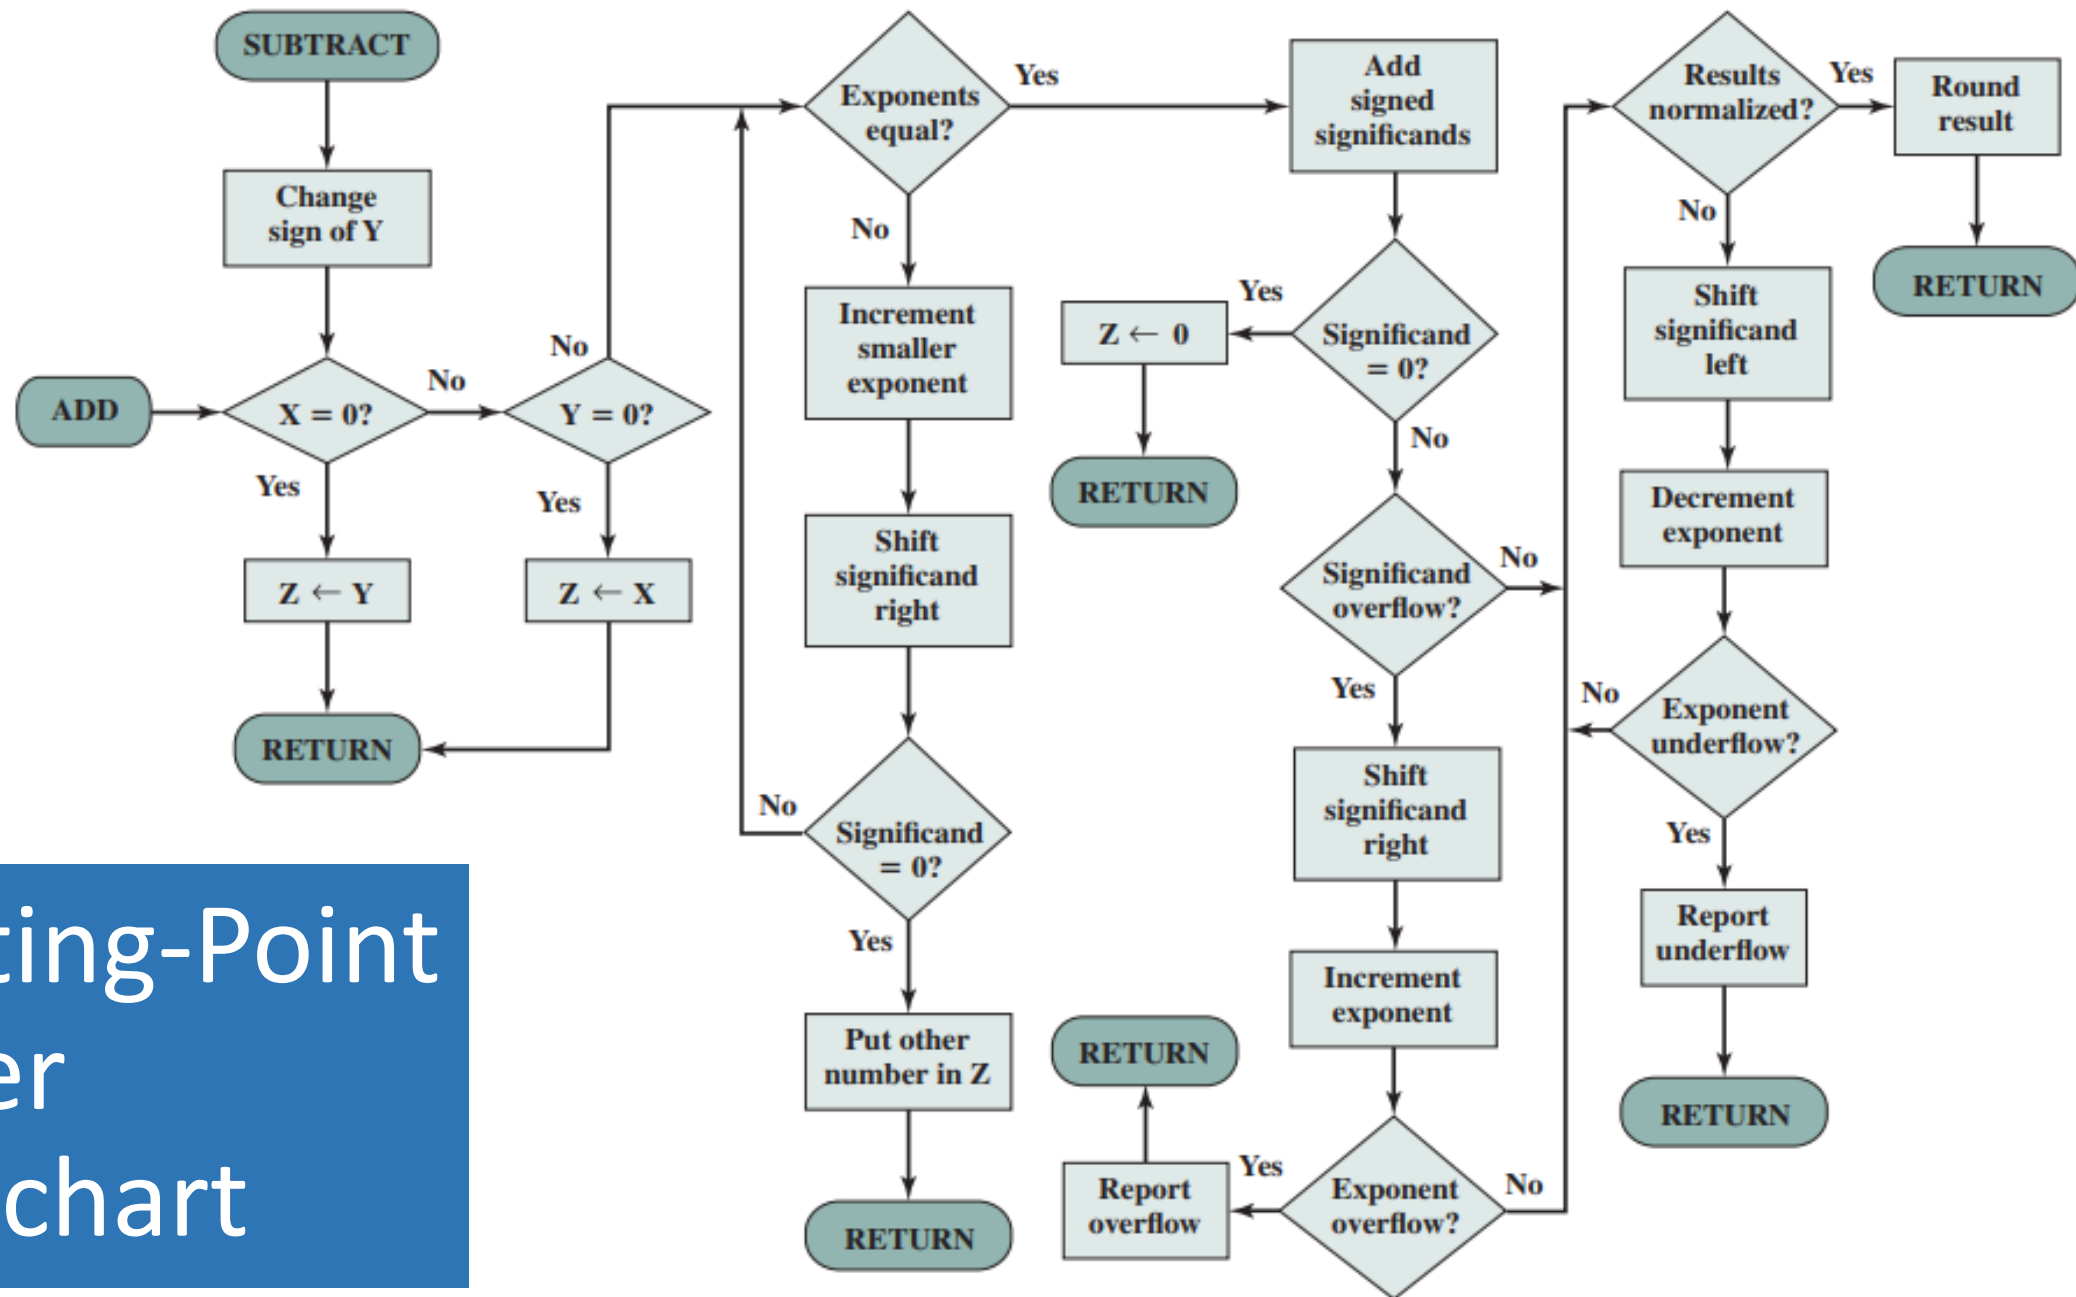
$$X \times Y = (X_S \times Y_S) \times B^{X_E + Y_E}$$

$$\frac{X}{Y} = \left(\frac{X_S}{Y_S} \right) \times B^{X_E - Y_E}$$

- Consider a 4-digit decimal example
 - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
 - Shift number with smaller exponent
 - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
 - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
 - 1.0015×10^2
- 4. Round and renormalize if necessary
 - 1.002×10^2

- Consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$ ($0.5 + -0.4375$)
- 1. Align binary points
 - Shift number with smaller exponent
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
 - $1.000_2 \times 2^{-4}$, with no over/underflow
- 4. Round and renormalize if necessary
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625

- Check for zeros
- Align significands (adjusting exponents)
- Add or subtract significands
- Normalize result



Floating-Point Adder Flowchart

- Chapter 3, P&H Textbook