

CS / EE 320

Computer Organization and Assembly Language

Lecture 25

Spring 2025

Shahid Masud

**Topics: Cache/Memory Update Write Strategies, Cache
Performance, Multi Level Caches, CD ROM Technology**

- Cache Writes – Write Back and Write Through schemes
- Examples of calculating different parameters in Cache Design
- Some Examples of Cache Performance
- Introduction to Multi-level Caches
- Examples of Cache Performance
- ---
- Maybe CD ROM Technology

Quiz 5 Next Lecture

Dealing with CPU to Write to Memory and Caches

- Must not overwrite a cache block unless main memory is up to date
- Multiple CPUs may have individual caches
- I/O may address main memory directly


Write-through scheme



- All writes go to main memory as well as cache
- Multiple CPUs can monitor main memory traffic to keep local (to CPU) cache up to date
- Lots of traffic
- Slows down writes

Write Through

- ⌘ Store by processor updates cache *and* memory
- ⌘ Memory always consistent with cache
- ⌘ ~2X more loads than stores
- ⌘ WT always combined with write buffers so that don't wait for lower level memory



- On data-write hit, could just update the block in cache
 - But then cache and memory would be inconsistent
- Write through: also update memory
- But makes writes take longer
 - e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
 - Effective CPI = $1 + 0.1 \times 100 = 11$
- Solution: write buffer
 - Holds data waiting to be written to memory
 - CPU continues immediately
 - Only stalls on write if write buffer is already full

- Updates initially made in cache only
- Update bit for cache slot is set when update occurs
- If block is to be replaced, write to main memory only if update bit is set
- Other caches get out of sync
- I/O must access main memory through cache
- N.B. 15% of memory references are writes

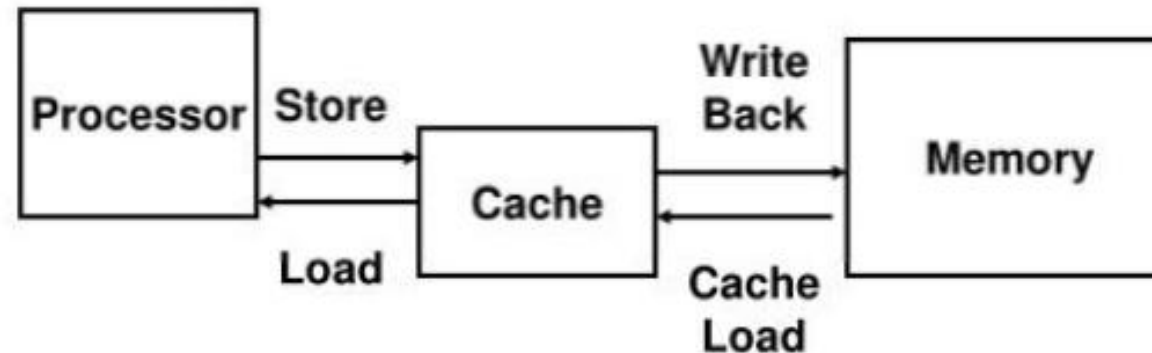
- Alternative: On data-write hit, just update the block in cache
 - Keep track of whether each block is dirty
- When a dirty block is replaced
 - Write it back to memory
 - Can use a write buffer to allow replacing block to be read first

Write Back Diagram

Write Back

Down

- ⌘ Store by processor only updates cache line
- ⌘ Modified line written to memory only when it is evicted
 - ☒ Requires "dirty bit" for each line
 - ☒ Set when line in cache is modified
 - ☒ Indicates that line in memory is stale
- ⌘ Memory not always consistent with cache
- ⌘ No writes of repeated writes



- What should happen on a write miss?
- Alternatives for write-through
 - Allocate on miss: fetch the block
 - Write around: don't fetch the block
 - Since programs often write a whole block before reading it (e.g., initialization)
- For write-back
 - Usually fetch the block

Design change	Effect on miss rate	Negative performance effect
Increase cache size	Decrease capacity misses	May increase access time
Increase associativity	Decrease conflict misses	May increase access time
Increase block size	Decrease compulsory misses	Increases miss penalty. For very large block size, may increase miss rate due to pollution.

How many memory references?

⌘ Each miss reads a block

☐ 2 bytes in this cache


⌘ Each store writes a byte

⌘ Total reads: 8 bytes

⌘ Total writes: 2 bytes

but caches generally miss $< 20\%$

- Example cache characteristics
 - Direct-mapped, write-back, write allocate
 - Block size: 4 words (16 bytes)
 - Cache size: 16 KB (1024 blocks)
 - 32-bit byte addresses
 - Valid bit and dirty bit per block
 - Blocking cache
 - CPU waits until access is complete



- **Multi Level Caches**
- **Performance of Cache and Memory System**
- **Examples of Cache Performance**

Example 1 of Cache Performance



Question

Given that:

T1 is access time from Cache = 0.01 μ sec

T2 is access time from Main Memory = 0.1 μ sec

95% accesses are from Cache

5% accesses are from Main Memory

Find the Average Time seen by CPU to Access one word?

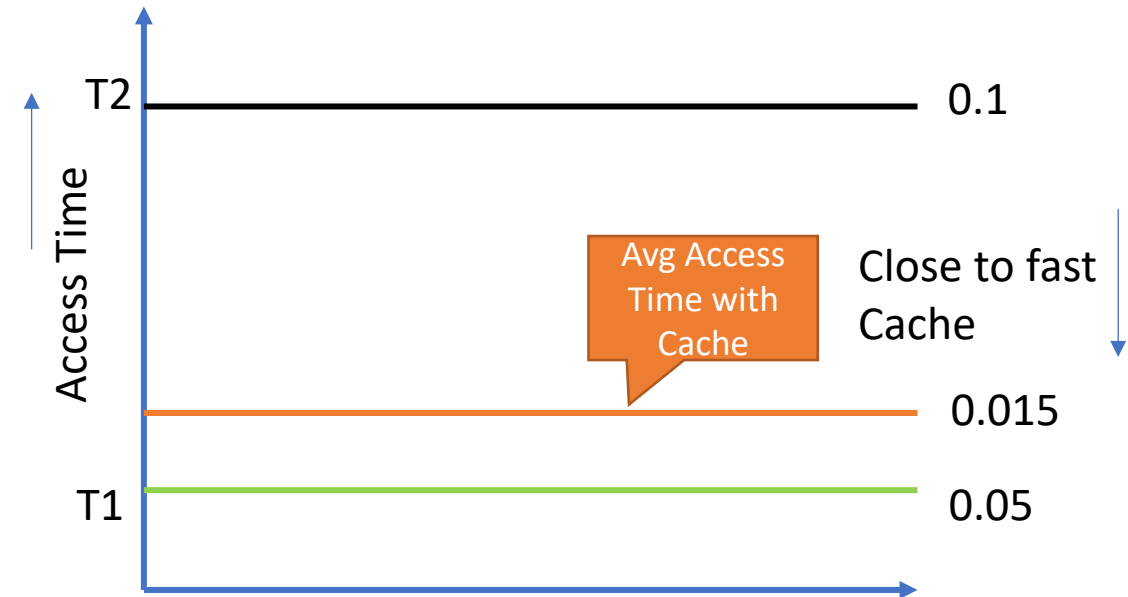
Solution

Average time to access one word:


$$T_{avg} = (0.95 \times 0.01) + (0.05 \times 0.1)$$

$$T_{avg} = 0.0095 + 0.0055$$

$$T_{avg} = 0.015 \mu\text{sec}$$




Separate Cache for Instructions and Data to take advantage of pipelining
In case the program has too many localized accesses to data only or instructions only,
Then there may be no performance improvement

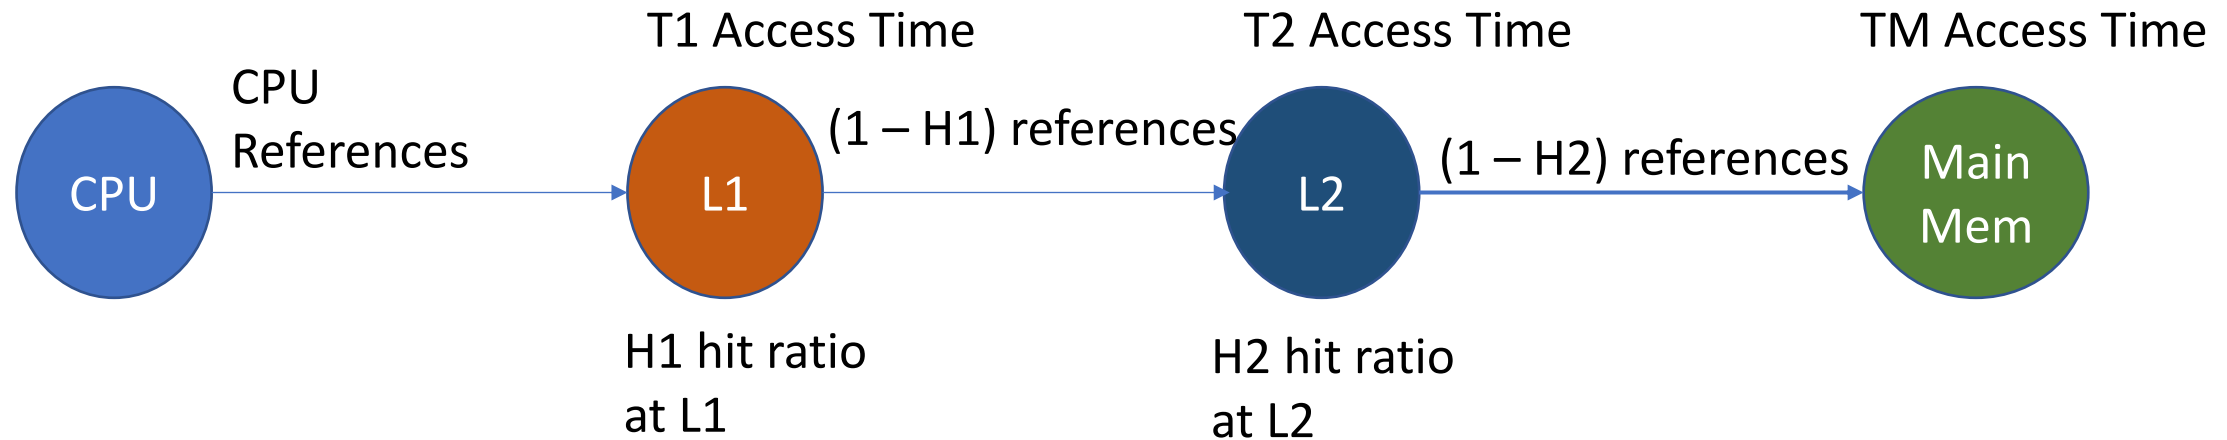


- Primary cache attached to CPU
 - Small, but fast
- Level-2 cache services misses from primary cache
 - Larger, slower, but still faster than main memory
- Main memory services L-2 cache misses
- Some high-end systems include L-3 cache

Main purpose is to reduce Miss Penalty



Multi Level Caches - Performance



Average Access Time in Multilevel Cache:

$$T_{avg} = T1 + (1 - H1) \times T2 + (1 - H1) \times (1 - H2) \times TM$$

$$(1 - H1) = \text{miss rate at L1} = \text{MRL1}$$

$$T_{avg} = T1 + \text{MRL1} \times T2 + \text{MRL1} \times \text{MRL2} \times TM$$

$$\text{MRL1} \times \text{MRL2} = \text{Global Miss Rate at L2}$$

Example 1, find Tavg in 2 level cache



- CPU generates 100 memory references. 80 are hit in L1, 20 are miss in L1. 16 are hit in L2 and 4 are miss in L2. Access time of L1 is 1 cycle. Access time of L2 is 10 cycles. Access time of Main Mem is 100 cycle. Find average memory access time Tavg?

- $T_{avg} = T_1 + MRL_1 \times T_2 + MRL_1 \times MRL_2 \times T_M$
- $= 1 + (20/100) \times 10 + [(20/100) \times (4/20)] \times 100$
- $= 1 + 2 + 4 = 6$ cycles Avg Access time

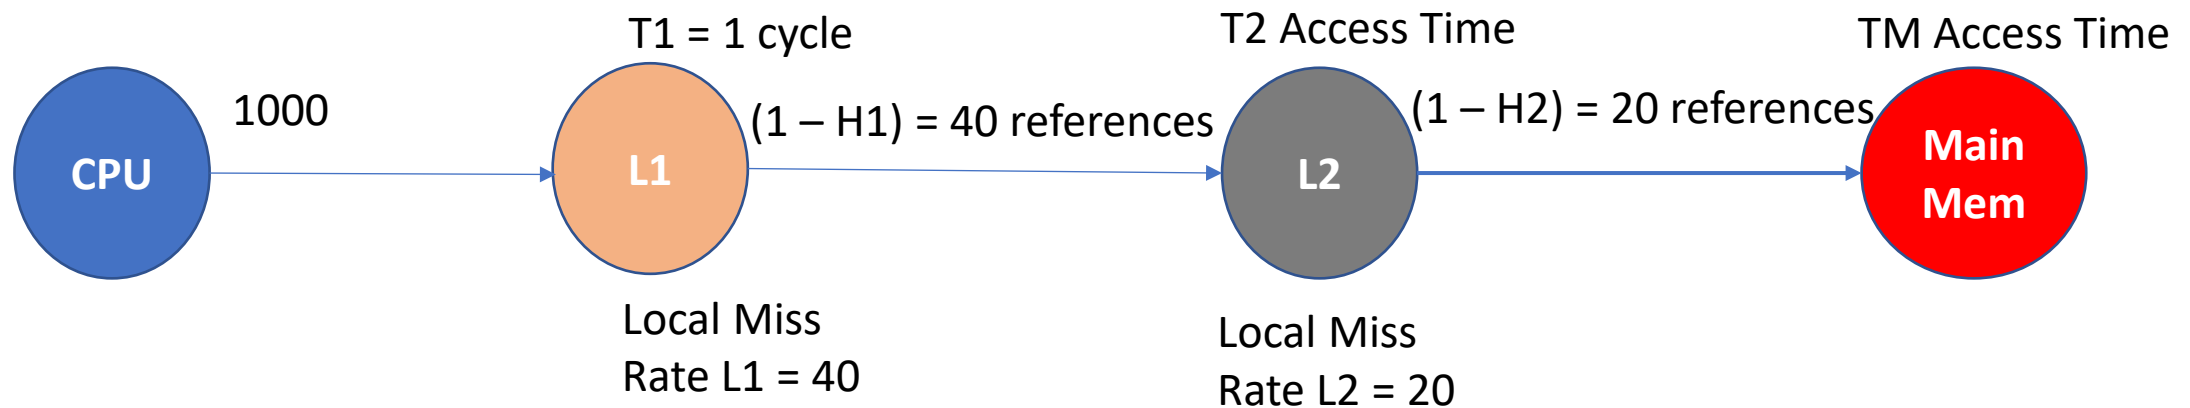
Note:

MRL1 = Miss Rate at Cache L1

Example 2, find Tavg in 2 level cache



- CPU generates 1000 memory references. Number of Miss in L1 is 40. Number of miss in L2 is 20. Access time of L1 is 1 cycle. Access time of L2 is 10 cycle. Find Tavg when Miss Penalty of L2 is 100 cycle.



$$\begin{aligned} T_{avg} &= T1 + MRL1 \times T2 + MRL1 \times MRL2 \times TM \\ &= 1 + (40/1000) \times 10 + [(40/1000) \times (20/40)] \times 100 \\ &= 1 + 0.4 + 2 = 3.4 \text{ cycles Avg Access time} \end{aligned}$$

- Components of CPU time
 - Program execution cycles
 - Includes cache hit time
 - Memory stall cycles
 - Mainly from cache misses
- With simplifying assumptions:

Memory stall cycles

$$= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty}$$

- Given
 - I-cache miss rate = 2%
 - D-cache miss rate = 4%
 - Miss penalty = 100 cycles
 - Base CPI (ideal cache) = 2
 - Load & stores are 36% of instructions
- Miss cycles per instruction
 - I-cache: $0.02 \times 100 = 2$
 - D-cache: $0.36 \times 0.04 \times 100 = 1.44$
- Actual CPI = $2 + 2 + 1.44 = 5.44$
 - Ideal CPU is $5.44/2 = 2.72$ times faster

- Hit time is also important for performance
- Average memory access time (AMAT)
 - $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- Example
 - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, l-cache miss rate = 5%
 - $AMAT = 1 + 0.05 \times 20 = 2\text{ns}$
 - 2 cycles per instruction

- When CPU performance increased
 - Miss penalty becomes more significant
- Decreasing base CPI
 - Greater proportion of time spent on memory stalls
- Increasing clock rate
 - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance


- Given
 - CPU base CPI = 1, clock rate = 4GHz
 - Miss rate/instruction = 2%
 - Main memory access time = 100ns
- With just primary cache
 - Miss penalty = $100\text{ns} / 0.25\text{ns} = 400$ cycles
 - Effective CPI = $1 + 0.02 \times 400 = 9$


Optical Drives CD / DVD

- Originally for audio
- 650Mbytes giving over 70 minutes audio
- Polycarbonate coated with highly reflective coat, usually aluminum
- Data stored as pits
- Read by reflecting laser
- Constant packing density
- Constant linear velocity

CD Components








Low-magnification ($\times 32$) image of a CD showing an edge of the data zone.

Spiral Track of CD




©2000 How Stuff Works


The CD is 12 cm in diameter, **1.2 mm thick**, has a center hole 1.5 cm in diameter, and spins at a *constant linear velocity* (CLV) or *constant angular velocity* (CAV).

There is only one track on the optical disk and all data are stored in a spiral of about **2 billion small pits** on the surface. There are about 30,000 windings on a CD - all part of the same track. This translates into about 16,000 tracks per inch and an areal density of 1 Mb/mm².

The total length of the track on a CD is almost 3 miles.

Track Density CD vs DVD



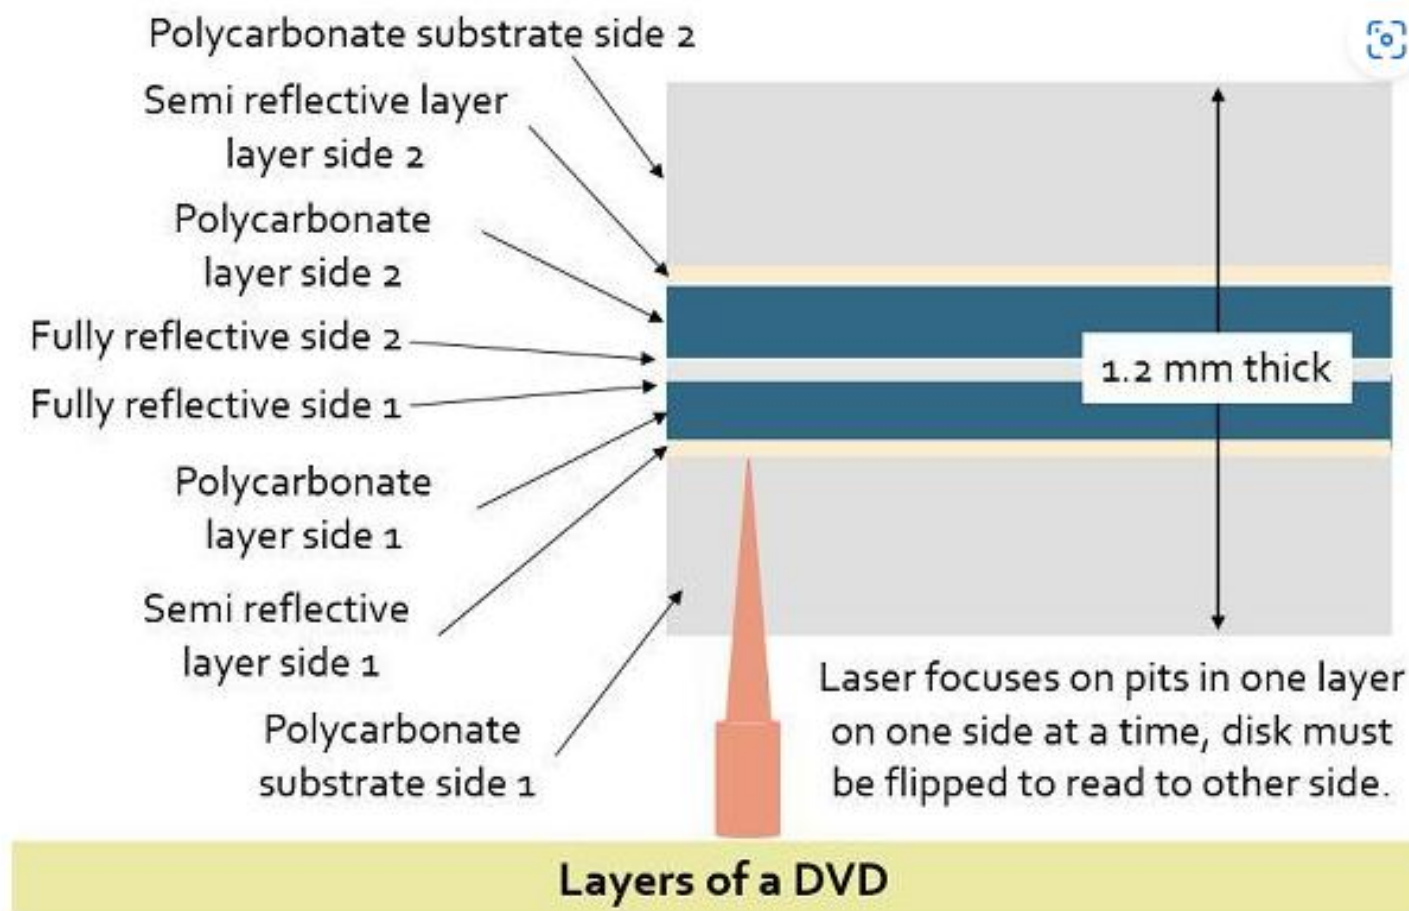



A CD can store up to 74 minutes of music, so the total amount of digital data that must be stored on a CD is:

$$2 \text{ channels} \times 44,100 \text{ samples/channel/second} \times 2 \text{ bytes/sample} \times 74 \text{ minutes} \times 60 \text{ seconds/minute} = 783,216,000 \text{ bytes}$$

To fit more than 783 megabytes onto a disk only 12 cm in diameter requires that the individual bits be very small.

Layers of a DVD





Digital data are carved into the disc as pits (low spots) and lands (high spots). As the laser shines into the moving pits and lands, a sensor detects a change in reflection when it encounters a transition from pit to land or land to pit. Each transition is a 1. The lack of transitions are 0s. There is only one laser in a drive. Two are used here to illustrate the difference in reflection.





Figure 6.9 CD Operation



001001001010000000010000010000101000001001000



Storage Write and Read Principle






Illustration: With a CD-R, binary information is stored as "burned" areas (0) and unburned areas (1) in the dye layer sandwiched between the protective polycarbonate and the reflective aluminum.

CD and DVD Read Details




(a) CD-ROM—Capacity 682 MB



(b) DVD-ROM, double-sided, dual-layer—Capacity 17 GB

- Audio is single speed
 - Constant linear velocity
 - 1.2 ms^{-1}
 - Track (spiral) is 5.27km long
 - Gives 4391 seconds = 73.2 minutes
 - Data 176.4 K bytes/s total capacity 774.57 M Bytes
- Other speeds are quoted as multiples
- e.g. 24x \sim 4 M Bytes/s (data transfer rate)
- The quoted figure is the maximum the drive can achieve



- Mode 0=blank data field
- Mode 1=2048 byte data+error correction
- Mode 2=2336 byte data

- Difficult
- Move head to rough position
- Set correct speed
- Read address
- Adjust to required location

- Chap 5 of P&H Textbook